Sai Srinivasa Karthik Varanasi[1]

[1]Department of Computer Science and Engineering (Cybersecurity), GITAM University

November 27, 2025

Posted on 27 Nov 2025 — CC-BY 4.0 — https://doi.org/10.36227/techrxiv.176425515.57150749/v1 — ePrints posted on TechRxiv are preliminary reports that are not peer reviewed. They should not b...

Post-Quantum Secure MQTT for IoT: Comprehensive Performance Evaluation Under Network Stress

Sai Srinivasa Karthik Varanasi
Department of Computer Science and Engineering (Cybersecurity)
GITAM University, Hyderabad 502329, Telangana, India
Email: varanasikarthik44@gmail.com, svaranas3@gitam.in

***Abstract*—The emergence of quantum computing poses an existential threat to current Internet of Things security infrastructure, which relies predominantly on RSA and elliptic curve cryptography for confidentiality and authentication. Shor's algorithm enables polynomial-time factorization that can break these classical cryptographic primitives, necessitating urgent migration to quantum-resistant alternatives. This paper presents the first comprehensive performance evaluation of six NIST-standardized and candidate post-quantum key encapsulation mechanisms (ML-KEM-512, ML-KEM-768, NTRU-Prime-hrss, BIKE-L1, HQC-128, Classic-McEliece-348864) integrated into the MQTT protocol across 4,608 experimental configurations spanning diverse real-world network conditions. We systematically evaluate connection establishment latency, message delivery reliability, decryption computational overhead, and throughput stability under systematically varied network latency (10-150 ms representing LAN to satellite links), packet loss rates (0-10% spanning excellent to degraded connectivity), payload sizes (128-1024 bytes covering sensor telemetry to multimedia streams), and message transmission rates (1-10 Hz from periodic monitoring to real-time control). Statistical analysis of 622,080 MQTT messages totaling 384 hours of simulation time reveals that NTRU-Prime-hrss achieves optimal performance-security balance with 98.1% delivery ratio and 26.1 ms median handshake latency across all tested conditions, while Classic-McEliece exhibits 278.3 ms handshake overhead rendering it impractical for resource-constrained deployments despite offering conservative security margins. We demonstrate through generalized linear modeling that transport-layer effects (TCP retransmissions, TLS record framing overhead) dominate post-quantum cryptographic computational costs, accounting for approximately 70% of end-to-end latency in high-loss scenarios. Our novel three-gate decision framework provides quantitative deployment criteria for IoT operators balancing security requirements against operational constraints. Complete implementation, experimental dataset, and analysis scripts available at https://github.com/vssk18/PQSM-RESEARCH.**

***Index Terms*—Post-quantum cryptography, MQTT protocol, Internet of Things security, key encapsulation mechanisms, network performance evaluation, lattice-based cryptography, code-based cryptography**

I. Introduction

THE rapid convergence of quantum computing advances and IoT ecosystem proliferation creates an urgent imperative for cryptographic migration in embedded systems. The global IoT device population exceeded 15 billion units in 2023 with projections approaching 75 billion by 2030, spanning critical infrastructure (smart grids, water treatment), healthcare (remote patient monitoring, medical devices), industrial automation (SCADA systems, robotics), and consumer applications (smart homes, wearables). These resource-constrained devices typically operate with 16-256 KB RAM, 32-512 KB flash storage, and 8-32 bit microcontrollers running at 8-120 MHz, creating significant challenges for cryptographic integration.

Shor's algorithm [1], discovered in 1994, provides polynomial-time quantum algorithms achieving $O((\log N)^3)$ complexity for integer factorization (breaking RSA) and discrete logarithm computation (breaking Diffie-Hellman and elliptic curve cryptography). While originally theoretical, recent quantum computing milestones have compressed threat timelines substantially. IBM demonstrated 433-qubit Osprey processors in 2022 [2], Google achieved quantum supremacy with 53-qubit Sycamore in 2019, and projections from leading quantum computing researchers indicate 20+ million noisy qubit systems capable of breaking RSA-2048 may emerge between 2030-2035 [3]. This timeline intersects with typical IoT device operational lifetimes (10-20 years for industrial sensors, 30+ years for smart grid infrastructure), creating immediate vulnerability windows.

MQTT (Message Queuing Telemetry Transport) has emerged as the dominant IoT messaging protocol, serving an estimated 75% of IoT deployments [4] and handling billions of daily publish-subscribe transactions across diverse application domains. Current secure MQTT implementations rely on TLS 1.2/1.3 with RSA-2048 or ECDH P-256 for session establishment, both vulnerable to Shor's algorithm once CRQCs (cryptographically-relevant quantum computers) emerge. The "harvest now, decrypt later" (HNDL) attack vector [5] presents immediate risks even before CRQC availability: adversaries can intercept and store encrypted IoT traffic today using classical recording infrastructure, then retrospectively decrypt accumulated ciphertexts once quantum computers become available. This threatens long-term confidential data including medical records (HIPAA mandates 50-year retention), critical infrastructure telemetry (decades-long operational data), and military communications (extended classification periods).

In response to quantum threats, NIST initiated a rigorous eight-year post-quantum cryptography standardization process in 2016, culminating in 2024 selections [6]. NIST standardized

ML-KEM (Module Lattice-Based Key Encapsulation Mechanism, formerly CRYSTALS-Kyber) as the primary quantum-resistant KEM, available in three security parameter sets: ML-KEM-512 (NIST Security Level 1, equivalent to AES-128), ML-KEM-768 (Level 3, equivalent to AES-192), and ML-KEM-1024 (Level 5, equivalent to AES-256). Additionally, NIST identified three Round 4 candidate KEMs for continued evaluation: BIKE (Bit Flipping Key Encapsulation), HQC (Hamming Quasi-Cyclic), and Classic-McEliece. These algorithms provide security based on hard mathematical problems believed intractable for both classical and quantum computers—specifically, lattice problems (Learning With Errors) for ML-KEM and NTRU variants, and syndrome decoding for code-based schemes.

However, integrating post-quantum cryptography into resource-constrained IoT environments presents substantial challenges beyond cryptographic security proofs. Post-quantum KEMs exhibit significantly larger cryptographic artifacts compared to classical schemes: public keys range from 800 bytes (ML-KEM-512) to 260 KB (Classic-McEliece-348864), ciphertexts span 768 bytes to 240 KB, and computational requirements increase 10-100× over ECDH operations. Network implications prove equally significant—large ciphertexts fragment across multiple IP packets (typical MTU 1500 bytes), potentially triggering retransmissions under packet loss and increasing connection establishment latency. Memory constraints on embedded devices may preclude certain algorithms entirely, while energy budgets for battery-powered sensors must accommodate increased cryptographic computation and transmission overhead.

A. Research Gap and Motivation

Existing post-quantum cryptography evaluations for IoT focus predominantly on isolated cryptographic microbenchmarks—measuring key generation time, encapsulation latency, and decapsulation overhead under idealized laboratory conditions with infinite memory, zero network latency, and perfect packet delivery. Representative examples include ARM Cortex-M4 benchmarking [7] reporting cycle counts for individual operations, and liboqs [8] performance characterization on desktop processors. While these measurements provide valuable insights into algorithmic computational complexity, they critically neglect system-level effects that dominate real-world IoT deployments.

Specifically, prior work exhibits four fundamental limitations. First, network impairment neglect: laboratory evaluations assume ideal conditions, failing to characterize performance degradation under realistic impairments. Wide-area IoT deployments experience 10-150 ms latency and 0.1-10% packet loss. Large post-quantum ciphertexts fragmenting across multiple packets exhibit nonlinear performance degradation when packets drop, as TCP retransmission timers trigger exponential backoff. Second, protocol stack overhead omission: isolated KEM benchmarks ignore multi-layer protocol processing. Real MQTT-over-TLS handshakes incur TCP three-way handshake, TLS record framing, cipher suite negotiation, certificate exchange, and MQTT protocol processing. Under high latency or loss, these protocol layers may dominate end-to-end performance, rendering cryptographic operation optimization marginally effective. Third, scalability blind spots: microbenchmarks measure single-connection performance, providing no insight into concurrent connection handling or denial-of-service vulnerability from slow handshakes. Production MQTT brokers must handle thousands of concurrent connections with frequent device reconnections. Fourth, deployment decision gap: existing work reports raw performance numbers without translating results into actionable deployment guidance.

This gap between isolated benchmarks and production requirements creates deployment uncertainty, potentially leading organizations to prematurely deploy KEMs that perform well in lab but fail operationally, over-provision infrastructure due to worst-case assumptions, or delay post-quantum migration awaiting better algorithms, leaving systems vulnerable to HNDL attacks.

B. Contributions

This paper bridges the benchmark-deployment gap through comprehensive end-to-end evaluation of post-quantum MQTT under realistic network stress, coupled with operator-oriented decision frameworks. Our specific contributions include:

Systematic Experimental Design at Unprecedented Scale: We execute 4,608 experimental configurations systematically exploring the full parameter space: 6 KEMs (ML-KEM-512, ML-KEM-768, NTRU-Prime-hrss, BIKE-L1, HQC-128, Classic-McEliece-348864) $\times$ 4 latencies (10, 50, 100, 150 ms) $\times$ 4 loss rates (0, 1, 5, 10%) $\times$ 4 payload sizes (128, 256, 512, 1024 bytes) $\times$ 4 message rates (1, 2, 5, 10 Hz) $\times$ 3 statistical replicates. Each configuration publishes 135 MQTT messages with QoS 1 (acknowledged delivery) over 135-second trials, totaling 622,080 messages analyzed across approximately 384 hours of simulation time. This scale exceeds prior work by 10-100×, enabling robust statistical inference and generalization across diverse deployment scenarios.

Rigorous Statistical Framework: We apply hypothesis testing (Kruskal-Wallis for equal medians with Bonferroni-corrected post-hoc Dunn's test), generalized linear models with logit link for delivery ratio prediction capturing latency-loss interaction effects, bootstrap confidence intervals for handshake latency distributions, Wilson score intervals for binomial delivery ratios, and quantile regression for tail latency characterization. All results report 95% confidence intervals, p-values, effect sizes, and goodness-of-fit metrics, enabling reproducibility.

Three-Gate Operator Decision Framework: We introduce a novel quantitative methodology translating performance metrics into deployment recommendations through three complementary gates: Gate 1 (Setup SLO Compliance) ensures connection establishment meets service-level objectives with $\mathbb{P}(T_h \leq 100 \text{ ms}) \geq 0.95$; Gate 2 (Reliability Envelope) requires robust delivery across network conditions with safe-area score $A_{\geq 0.95} \geq 0.90$; Gate 3 (Throughput Elasticity) bounds goodput degradation under packet loss with $|\varepsilon_{G,p}| \leq 2.0$.

Performance Decomposition Revealing Transport Dominance: Through detailed latency instrumentation using eBPF

probes capturing sub-microsecond timestamps across protocol layers, we demonstrate that transport effects (TCP retransmissions 42%, TLS overhead 28%) dominate cryptographic computational costs (KEM operations 14%) in wide-area high-loss scenarios. This finding challenges conventional optimization focus on accelerating isolated KEM operations, suggesting network stack tuning yields 2-3× greater performance gains.

Actionable Deployment Findings: Our evaluation establishes that NTRU-Prime-hrss achieves optimal overall balance (98.1% delivery ratio, 26.1 ms median handshake, perfect 1.00 safe-area score). ML-KEM-512 offers fastest handshakes (21.9 ms median) for latency-critical applications. ML-KEM-768 provides enhanced security with acceptable performance for high-assurance systems. Critically, Classic-McEliece fails practical deployment requirements (278.3 ms median handshake, 0.94 safe-area score) despite offering most conservative security assumptions.

Open Reproducible Research Artifact: We release complete implementation (3,200+ lines spanning KEM integration, MQTT broker modifications, network emulation infrastructure, eBPF monitoring probes, and statistical analysis), full experimental dataset (622,080 message records with per-packet timestamps), and analysis notebooks at https://github.com/vssk18/PQSM-RESEARCH.

The remainder of this paper proceeds as follows. Section II provides background on MQTT protocol architecture, post-quantum cryptographic mathematical foundations, NIST security levels, and reviews related work. Section III details our framework architecture, protocol extensions for KEM integration into TLS 1.3, and implementation. Section IV presents experimental methodology including full factorial design, metric definitions, statistical methods, and three-gate decision framework formulation. Section V reports comprehensive results. Section VI discusses performance-security trade-offs, transport dominance analysis, phased migration strategy, and study limitations. Section VII concludes with key findings.

II. Background and Related Work

A. MQTT Protocol Architecture

Message Queuing Telemetry Transport represents a lightweight publish-subscribe messaging protocol specifically designed for constrained devices and unreliable networks. Originally developed by IBM and Eurotech in 1999 for supervisory control and data acquisition systems monitoring oil pipelines via satellite links, MQTT has evolved into an OASIS standard (v3.1.1 in 2014, v5.0 in 2019) now serving as the de facto IoT messaging backbone.

MQTT employs a broker-centric architecture fundamentally different from request-response protocols. The MQTT broker serves as central message hub, accepting TCP connections from clients, maintaining topic subscription registries, routing published messages to matching subscribers based on topic filters supporting wildcards, managing Quality-of-Service guarantees through message persistence and acknowledgment tracking, and handling client session state across disconnections.

MQTT clients perform two primary operations. PUBLISH operations send application messages to specified topics with optional QoS level and retain flag. SUBSCRIBE operations register interest in topic patterns using hierarchical wildcards: single-level wildcard matches one hierarchy level, while multi-level wildcard matches any number of trailing levels.

MQTT defines three QoS levels balancing reliability against overhead. QoS 0 (At Most Once) provides fire-and-forget delivery with no acknowledgments, suitable for high-frequency telemetry tolerating occasional loss. QoS 1 (At Least Once) provides acknowledged delivery ensuring messages arrive at least once but potentially duplicated, suitable for state updates and commands. QoS 2 (Exactly Once) provides four-way handshake guaranteeing exactly-once delivery without duplicates, suitable for billing transactions and critical commands. Our evaluation focuses on QoS 1 as it represents the dominant production use case.

Standard MQTT connection establishment proceeds through layered handshakes. First, client initiates TCP three-way handshake establishing reliable byte stream, consuming 1.5 round-trip times. Second, TLS handshake establishes encrypted channel: ClientHello advertises supported cipher suites, ServerHello selects cipher and provides certificate, client verifies certificate and performs key exchange, both sides derive session keys, consuming 1-2 RTT. Third, client sends MQTT CONNECT packet and broker responds with CONNACK. The TLS handshake phase—specifically the key exchange operation—represents the primary integration point for post-quantum cryptography.

B. Post-Quantum Cryptography Foundations

1) Mathematical Problem Hardness: NIST post-quantum KEMs derive security from hard mathematical problems believed intractable for both classical and quantum computers, unlike RSA and ECDH which succumb to Shor's algorithm.

Lattice-based cryptography (ML-KEM, NTRU) relies on Learning With Errors, a problem involving distinguishing noisy linear equations from random:

$$\mathbf{b} = \mathbf{A}\mathbf{s} + \mathbf{e} \pmod{q} \tag{1}$$

where $\mathbf{A} \in \mathbb{Z}_q^{m \times n}$ is uniformly random public matrix, $\mathbf{s} \in \mathbb{Z}_q^n$ is secret key vector drawn from narrow discrete Gaussian distribution, $\mathbf{e} \in \mathbb{Z}_q^m$ is small error vector, and q is prime modulus. Given $(\mathbf{A}, \mathbf{b})$, adversaries must recover $\mathbf{s}$ or distinguish $\mathbf{b}$ from uniform random. Best known quantum attacks use lattice reduction achieving $2^{0.292\sqrt{n}}$ operations [9], compared to classical $2^{0.292n}$—quantum speedup but still exponential security.

Code-based cryptography (BIKE, HQC, Classic-McEliece) relies on syndrome decoding for linear error-correcting codes:

$$\mathbf{s} = \mathbf{H}\mathbf{e}^T \tag{2}$$

where $\mathbf{H} \in \mathbb{F}_2^{r \times n}$ is parity-check matrix defining the code, $\mathbf{e} \in \mathbb{F}_2^n$ is error vector with Hamming weight w, and $\mathbf{s} \in \mathbb{F}_2^r$ is syndrome. Given $(\mathbf{H}, \mathbf{s})$, adversaries must find $\mathbf{e}$ with weight w satisfying the equation. Quantum speedup via Grover's algorithm achieves $2^{0.5w}$ complexity [10], compared to classical 2^w, but exponential security remains.

Classic-McEliece uses structured Goppa codes offering most conservative security with no structural attacks in 40+ years. BIKE and HQC use quasi-cyclic codes enabling smaller keys but facing ongoing cryptanalytic evaluation.

2) NIST Security Levels: NIST categorizes post-quantum algorithms by quantum attack computational cost relative to AES key search. Level 1 requires cost equivalent to AES-128 key recovery ($\geq 2^{143}$ quantum gates), suitable for commercial applications requiring 10-20 year confidentiality. Level 3 requires cost equivalent to AES-192 ($\geq 2^{207}$ gates), suitable for government and high-security applications requiring 30-50 year confidentiality. Level 5 requires cost equivalent to AES-256 ($\geq 2^{272}$ gates), appropriate for military and ultra-long-term confidentiality.

Parameter selection involves balancing security margins against performance and size overhead. ML-KEM-512 (Level 1) offers 800-byte public keys and achieves 21.9 ms handshakes in our evaluation, while ML-KEM-1024 (Level 5) incurs 1568-byte keys with estimated 40+ ms handshakes. For most IoT applications, Level 1 provides adequate protection against foreseeable quantum threats, with Level 3 recommended for critical infrastructure.

C. Related Work

1) Post-Quantum Cryptography for Embedded Systems: Kannwischer et al. [7] introduced pqm4, a comprehensive benchmarking framework for NIST Round 2 post-quantum candidates on ARM Cortex-M4 microcontrollers. Their evaluation reported cycle counts: Kyber-512 required 830K cycles for key generation, 920K for encapsulation, 930K for decapsulation on 168 MHz STM32F4. While providing valuable algorithmic complexity insights, pqm4 evaluations measure isolated operations without network integration or statistical confidence intervals. Crockett et al. [13] extended this work with optimized implementations achieving 2.5× speedup through assembly-level NEON vectorization on ARM Cortex-A processors, though still focusing on isolated KEM operations. Fritzmann et al. [14] evaluated CRYSTALS-Kyber on ESP32 microcontrollers commonly deployed in IoT, reporting 15-20 ms key generation times but neglecting protocol integration overhead.

Alkim et al. [15] analyzed memory requirements for NIST PQC candidates, identifying Classic-McEliece's 260 KB public keys as incompatible with resource-constrained devices having less than 512 KB RAM. Howe et al. [16] proposed hardware acceleration for lattice-based schemes using FPGA implementations, achieving 10× performance improvements but requiring custom silicon impractical for commercial IoT deployments. Recent work by Saarinen [17] demonstrated constant-time implementations preventing side-channel attacks, critical for production deployment but adding 15-20% computational overhead.

2) Post-Quantum TLS Integration: Sikeridis et al. [11] conducted measurements of post-quantum TLS 1.3 across diverse client platforms, finding that lattice-based schemes imposed 2-3× handshake overhead compared to ECDH under low-latency LAN conditions (RTT less than 1 ms). However, their experiments did not characterize performance under packet loss or high-latency scenarios. Schwabe et al. [18] evaluated hybrid TLS combining classical ECDH with post-quantum KEMs, demonstrating backward compatibility but doubling handshake overhead. Paquin et al. [19] analyzed real-world TLS 1.3 deployments at Cloudflare, finding that 2 KB ciphertext size threshold determines whether fragmentation impacts user experience, though this study excluded IoT scenarios.

Stebila and Mosca [20] provided early analysis of post-quantum TLS integration challenges, identifying large ciphertext sizes as primary deployment barrier. Recent work by Kampanakis and Panburana [21] measured hybrid TLS 1.3 performance across 50 geographic regions, finding that ML-KEM-768 adds 15-25 ms median latency under realistic internet conditions. Kwiatkowski and Valenta [22] conducted longitudinal study of Chrome browser TLS handshakes with post-quantum algorithms, observing 98.5% success rate but 4% increase in connection timeouts.

3) MQTT Security and IoT Protocols: Andy et al. [23] analyzed MQTT security vulnerabilities, identifying weak authentication and lack of message integrity in unencrypted deployments. Mishra et al. [24] proposed lightweight authentication schemes for MQTT reducing handshake overhead by 40% through session resumption, though still relying on vulnerable RSA cryptography. Recent work by Naik [25] evaluated MQTT over QUIC transport, achieving 30% latency reduction compared to TCP but not addressing post-quantum threats.

Calabretta et al. [12] explored post-quantum cryptography integration into MQTT, implementing Kyber-based key exchange in simulated broker under ideal network conditions. Their evaluation measured connection establishment time under zero packet loss, finding acceptable overhead for Kyber-512 and Kyber-768. However, their work lacked systematic stress testing, statistical rigor, and scalability evaluation. Porambage et al. [26] surveyed IoT security protocols, highlighting MQTT's dominance but noting inadequate protection against quantum threats. Vegh and Miclea [27] proposed edge-based MQTT security architecture distributing cryptographic computation, reducing device-side overhead by 60% but increasing infrastructure complexity.

4) Network Performance Evaluation Methodologies: Hemminger [28] introduced Network Emulation (netem) enabling controlled packet loss and latency injection at kernel level, becoming standard for reproducible network research. Carbone and Rizzo [29] validated netem accuracy against physical testbeds, finding less than 2% deviation for latency emulation and 5% for loss models. Recent extensions by Leclaire [30] added time-varying conditions supporting mobile IoT scenarios.

Botta et al. [31] surveyed IoT traffic characterization, establishing that real-world deployments experience 10-150 ms latency and 0.1-10% packet loss depending on technology (WiFi, 4G, satellite). Augustin et al. [32] measured TCP performance under packet loss, demonstrating exponential backoff amplification when loss exceeds 5%. Langley et al. [33] analyzed QUIC transport protocol resilience to packet loss, motivating similar analysis for MQTT over TCP.

5) Statistical Analysis in Performance Evaluation: Jain [34] established experimental design principles for computer systems, advocating full factorial designs for characterizing interaction effects. Georges et al. [35] demonstrated importance of statistical rigor in benchmarking, showing that single-run results often mislead due to variance. Kalibera and Jones [36] provided guidance on confidence interval estimation for performance measurements, recommending bootstrap methods for non-normal distributions.

Dunn [37] introduced rank-based post-hoc tests for non-parametric hypothesis testing, enabling pairwise comparisons while controlling family-wise error rate. McCullagh and Nelder [38] formalized generalized linear models for bounded response variables, applicable to delivery ratio prediction. Recent work by Chen et al. [39] applied GLM to network performance analysis, demonstrating superior fit compared to linear regression for reliability modeling.

6) IoT Security Threats and Migration Strategies: Fernandez et al. [40] analyzed harvest-now-decrypt-later threats against IoT, estimating that adversaries have captured 5+ years of encrypted traffic awaiting quantum decryption. Mosca and Piani [5] updated quantum threat timelines based on recent hardware advances, compressing CRQC availability projections to 2030-2035. Bindel et al. [41] proposed phased migration strategies balancing immediate quantum resistance against operational disruption, recommending hybrid classical-PQC deployments during transition periods.

Ghosh et al. [42] analyzed IoT device lifetime implications for cryptographic agility, finding that 30-year smart grid sensors must receive over-the-air firmware updates for post-quantum migration. Kumar et al. [43] measured energy consumption of post-quantum operations on battery-powered sensors, reporting that ML-KEM-512 adds 8-12% to connection establishment energy. Menezes and Campagna [44] discussed practical challenges in post-quantum migration for resource-constrained devices, identifying memory limitations as primary barrier.

7) Research Gap: Existing work exhibits four fundamental limitations. First, PQC benchmarking studies focus on isolated cryptographic operations without protocol integration, neglecting multi-layer overhead from TCP, TLS, and application protocols. Second, TLS performance evaluations assume ideal network conditions, failing to characterize degradation under packet loss and latency combinations typical of wide-area IoT. Third, MQTT security research proposes post-quantum integration but lacks statistical rigor, systematic stress testing, or scalability analysis. Fourth, performance studies report raw metrics without translating results into actionable deployment guidance through quantitative decision frameworks.

We provide the first comprehensive end-to-end evaluation bridging cryptographic primitives, network protocols, and operational decision-making through systematic experimental design at 10-100× larger scale than prior work, realistic stress testing across latency-loss parameter space, rigorous statistics with hypothesis testing and confidence bounds, quantitative three-gate decision framework, performance decomposition revealing transport dominance, concurrent scalability evaluation, and open reproducible artifact.

III. System Design and Implementation

A. Framework Architecture

Our Post-Quantum Secure Messaging framework extends standard MQTT with quantum-resistant key exchange while maintaining backward compatibility and preserving MQTT protocol semantics. The architecture comprises four primary components.

First, PQC-enabled MQTT clients modified from Eclipse Paho Python client library support post-quantum KEMs through TLS callback hooks. Clients implement cipher suite negotiation, advertising supported KEMs in TLS extensions, hybrid classical-PQC modes for gradual migration, and graceful fallback to classical cryptography when connecting to legacy brokers. Client-side modifications totaled approximately 800 lines of Python integrating liboqs bindings for KEM operations.

Second, PQC-enabled MQTT broker extended from Eclipse Mosquitto 2.0 with custom authentication callbacks interfacing with liboqs for KEM operations. Broker maintains heterogeneous connection state, serving concurrent classical and post-quantum clients simultaneously during migration periods, managing per-client cipher suite selections, enforcing QoS guarantees across PQC and classical sessions, and persisting session state to disk for recovery after crashes. Broker modifications spanned approximately 1,200 lines of C integrating liboqs and extending authentication subsystem.

Third, network impairment simulator leverages Linux Traffic Control with Network Emulation kernel module to inject controlled latency, packet loss, bandwidth limits, and jitter at packet level, providing realistic wide-area network simulation without requiring physical testbed infrastructure. The netem subsystem operates at the link layer, buffering outgoing packets and applying stochastic delays and drops before transmission, accurately modeling correlated loss bursts and latency variance observed in production networks. Configuration scripts (approximately 400 lines Bash) systematically apply 256 network condition combinations across experimental runs.

Fourth, monitoring and analysis layer deploys eBPF probes to capture sub-microsecond timestamps of critical events: TCP connection state transitions, TLS handshake message boundaries, KEM operation entry and exit, and MQTT packet timestamps. eBPF enables tracing with near-zero overhead (less than 1% CPU, less than 100 KB memory) by executing sandboxed bytecode directly in kernel space, avoiding expensive context switches to userspace. Monitoring infrastructure comprises approximately 800 lines of C (eBPF probes) and 600 lines of Python (event aggregation and analysis).

B. Post-Quantum Protocol Extensions

We integrate KEMs into TLS 1.3 handshake by replacing the KeyShare extension traditionally containing ECDH public keys. The modified handshake proceeds through six phases.

Phase 1 (TCP Establishment): Client initiates TCP three-way handshake to broker port 8883, completing in 1.5 RTT consuming $T_{\text{TCP}} = 1.5 \times \text{RTT}$.

Phase 2 (TLS Hello): Client sends ClientHello containing supported cipher suites and PQC KEM algorithms in

extensions field. Broker responds with ServerHello selecting KEM algorithm (e.g., ML-KEM-512) and including server's ephemeral KEM public key pk generated via $(pk, sk) \leftarrow \text{KeyGen}()$.

Phase 3 (Key Exchange): Client performs KEM encapsulation $(ct, ss) \leftarrow \text{Encaps}(pk)$ generating random shared secret ss and its encapsulation (ciphertext) ct. Client sends ClientKeyExchange message containing ct (768 bytes for ML-KEM-512, potentially fragmenting across multiple TCP segments for larger KEMs).

Phase 4 (Server Decapsulation): Broker receives ciphertext ct and performs decapsulation $ss' \leftarrow \text{Decaps}(ct, sk)$ recovering the shared secret. Mathematical correctness guarantees $ss' = ss$ with overwhelming probability.

Phase 5 (Key Derivation): Both client and broker derive session keys using HKDF (HMAC-based Key Derivation Function):

$$k_{\text{session}} = \text{HKDF}(ss, \text{nonce}_{\text{client}} \| \text{nonce}_{\text{server}}, \text{label}) \quad (3)$$

generating 128-bit AES-GCM encryption key and authentication key.

Phase 6 (MQTT Protocol): Client sends encrypted MQTT CONNECT packet containing client identifier, username and password credentials, clean session flag, and keepalive interval. Broker validates credentials, allocates session resources, and responds with encrypted CONNACK packet indicating success or failure. Connection is now established for subsequent PUBLISH and SUBSCRIBE operations.

Total handshake time measured from TCP SYN to MQTT CONNACK receipt:

$$T_h = T_{\text{TCP}} + T_{\text{TLS-hello}} + T_{\text{encaps}} + T_{\text{network}}(ct) + T_{\text{decaps}} + T_{\text{derive}} + T_{\text{MQTT}} \quad (4)$$

For defense-in-depth, we implement hybrid key exchange combining ECDH X25519 with post-quantum KEMs, deriving session key as:

$$k_{\text{hybrid}} = \text{HKDF}(ss_{\text{PQC}} \| ss_{\text{ECDH}}, \text{nonces}, \text{label}) \quad (5)$$

Security holds if either component remains unbroken, providing protection against potential post-quantum cryptanalytic breakthroughs or undiscovered quantum algorithm improvements.

C. Implementation Details

We utilize Open Quantum Safe liboqs 0.9.0, providing optimized implementations of NIST-standardized and candidate algorithms. liboqs incorporates assembly-optimized code paths for x86-64 AVX2 instructions, achieving 2-3× speedup over portable C implementations for lattice-based schemes.

Post-KEM session encryption uses AES-128-GCM (Galois/Counter Mode) for authenticated encryption, providing confidentiality and integrity with 128-bit security. AES-GCM benefits from hardware acceleration via AES-NI instructions on modern x86 processors, achieving multi-gigabit throughput. We configure 12-byte nonces and 16-byte authentication tags per NIST recommendations.

Broker configured with QoS 1 default, 60-second keepalive interval, 30-second message timeout, and 10,000-client connection limit. We disable QoS 2 and retained messages to isolate handshake and delivery performance. Persistence enabled to disk for session recovery. Clients configured with automatic reconnect on connection loss, exponential backoff starting at 1 second up to 60 seconds maximum.

Network emulation via tc netem configuration applies per-interface to broker ingress, affecting all incoming client connections uniformly. Latency configured with normal distribution (mean L, standard deviation $0.1L$) modeling jitter. Loss configured with independent uniform random model at specified rate p. We verified netem accuracy through independent validation measuring actual end-to-end RTT and loss rate against configured values, observing less than 2% deviation.

Hardware platform: broker deployed on Intel Xeon Gold 6248R at 3.0 GHz (24 cores, 48 threads), 128 GB DDR4-2933 ECC RAM, Intel X710 10 Gbps Ethernet, Samsung 970 PRO 1 TB NVMe SSD, Ubuntu 22.04 LTS (kernel 5.15). Clients distributed across 8 physical hosts: AMD Ryzen 9 5900X at 3.7 GHz (12 cores, 24 threads), 64 GB DDR4-3200 RAM, Realtek RTL8125B 2.5 Gbps Ethernet, Ubuntu 22.04 LTS. All systems connected via dedicated Gigabit Ethernet switch with jumbo frames disabled (MTU 1500), isolated VLAN.

IV. Experimental Methodology

A. Systematic Experimental Design

We employ full factorial experimental design systematically exploring all combinations of factors to characterize main effects and interaction effects. Five primary factors span the IoT deployment space.

KEM Algorithm: 6 algorithms (ML-KEM-512, ML-KEM-768, NTRU-Prime-hrss, BIKE-L1, HQC-128, Classic-McEliece-348864) covering NIST standardized and candidates from lattice-based and code-based families.

Network Latency: 4 levels (10, 50, 100, 150 ms) representing LAN (10 ms), WiFi and 4G (50 ms), cellular backhaul (100 ms), and satellite (150 ms).

Packet Loss: 4 levels (0, 1, 5, 10%) representing excellent (0%), good (1%), degraded (5%), and poor (10%) connectivity.

Payload Size: 4 levels (128, 256, 512, 1024 bytes) representing sensor telemetry (128 B), aggregated data (256 B), rich telemetry (512 B), and multimedia stream (1024 B).

Message Rate: 4 levels (1, 2, 5, 10 Hz) representing periodic monitoring (1 Hz), event-driven (2-5 Hz), and real-time control (10 Hz).

Statistical Replicates: 3 independent runs to quantify measurement variance and enable confidence interval estimation.

Total configurations: $6 \times 4 \times 4 \times 4 \times 4 \times 3 = 4,608$ experimental runs.

Each experimental run proceeds as follows. First, configure broker with selected KEM and apply network emulation parameters (L, p). Second, start monitoring subsystem (eBPF probes, packet capture). Third, client establishes MQTT connection, measuring handshake time T_h. Fourth, client publishes 135 messages at rate f Hz with payload size B bytes, topic "test/performance", QoS 1. Fifth, record delivery

outcomes (PUBACK received or 30s timeout), decrypt latencies, timestamps. Sixth, client disconnects gracefully. Seventh, aggregate metrics, clean state, proceed to next configuration.

Total experiment duration: 135 messages times (1 divided by f) seconds per run times 4,608 runs approximately equals 384 hours pure measurement time, plus approximately 50 hours configuration overhead, totaling 18 days wall-clock time executing across 8 client hosts in parallel.

B. Performance Metrics

Handshake time measures connection establishment latency from TCP SYN transmission to MQTT CONNACK receipt, captured via eBPF timestamps:

$$T_h = t_{\text{CONNACK}} - t_{\text{SYN}} \text{ [milliseconds]} \tag{6}$$

Includes all protocol layers: TCP 3-way handshake (1.5 RTT), TLS negotiation with KEM operations, MQTT CONNECT and CONNACK processing. We report median (P50) as primary statistic (robust to outliers) and 95th percentile (P95) characterizing tail latency. Bootstrap 95% confidence intervals computed via 1000 resamples with replacement.

Delivery ratio measures fraction of published messages successfully acknowledged within 30-second timeout:

$$D = \frac{n_{\text{ack}}}{n_{\text{sent}}} = \frac{n_{\text{ack}}}{135} \tag{7}$$

QoS 1 messages count as delivered upon receiving PUBACK from broker. Messages failing to receive PUBACK within 30s count as lost. $D = 1.0$ indicates perfect delivery; $D = 0.95$ indicates 5% loss; $D < 0.90$ considered unacceptable for production systems. We report Wilson score confidence interval for binomial proportion (more accurate than normal approximation for proportions near 0 or 1).

Decrypt latency measures time from encrypted payload receipt to plaintext availability at broker:

$$t_{\text{dec}} = t_{\text{plaintext}} - t_{\text{ciphertext}} \text{ [milliseconds]} \tag{8}$$

Captures KEM decapsulation overhead plus symmetric decryption (AES-GCM). We report P50 and P95 across all messages in a run, then aggregate statistics across runs. Power-law scaling model:

$$t_{\text{dec}} = \alpha + \beta B^{\gamma} \tag{9}$$

where α captures fixed KEM overhead, βB^{γ} represents symmetric decryption scaling with payload B.

Goodput measures effective application-layer throughput accounting for message delivery failures:

$$G = \frac{f \cdot B \cdot 8 \cdot D}{10^6} \text{ [Megabits per second]} \tag{10}$$

where f is message rate [Hz], B is payload [bytes], 8 converts bytes to bits, D is delivery ratio. Goodput represents useful data rate from application perspective, discounting overhead from headers, retransmissions, and lost messages.

C. Statistical Analysis Methods

We test whether handshake time distributions differ significantly across KEMs using non-parametric Kruskal-Wallis test (appropriate for non-normal distributions, unequal variances):

$$H = \frac{12}{N(N+1)} \sum_{i=1}^{k} \frac{R_i^2}{n_i} - 3(N+1) \sim \chi^2_{k-1} \tag{11}$$

where $N = \sum_{i=1}^{k} n_i$ is total sample size, R_i is sum of ranks for group i, n_i is group i size, $k = 6$ KEMs. Null hypothesis: all KEMs have equal median handshake time. Rejection ($p < 0.001$) indicates significant difference exists.

Post-hoc pairwise comparisons use Dunn's test with Bonferroni correction controlling family-wise error rate: $\alpha' = \alpha/\binom{k}{2} = 0.05/15 = 0.00333$. Each pairwise test computes z-statistic from rank differences. Significant pairwise difference declared when absolute z-value exceeds Bonferroni-corrected critical value.

Bootstrap percentile method provides distribution-free 95% confidence interval for handshake time. Algorithm: resample observations with replacement from original dataset, compute median of resample, repeat 1000 times obtaining bootstrap distribution of medians, report 2.5th and 97.5th percentiles as CI bounds.

Wilson score interval for binomial proportion provides confidence interval for delivery ratio:

$$\text{CI} = \frac{\hat{p} + \frac{z^2}{2n} \pm z\sqrt{\frac{\hat{p}(1-\hat{p})}{n} + \frac{z^2}{4n^2}}}{1 + \frac{z^2}{n}} \tag{12}$$

where $\hat{p} = n_{\text{ack}}/135$ is observed proportion, $n = 135$ is sample size, $z = 1.96$ for 95% CI.

We model delivery ratio as function of network conditions using GLM with logit link (appropriate for bounded response):

$$\text{logit}(E[D|L,p]) = \log\left(\frac{E[D]}{1-E[D]}\right) = \beta_0 + \beta_L L + \beta_p p + \beta_{Lp}(L \times p) \tag{13}$$

Interaction term β_{Lp} captures non-additive joint effect: latency and loss together degrade delivery more than sum of individual effects due to exponential TCP backoff amplification. Model fit via maximum likelihood estimation.

Safe-area score quantifies reliability robustness as fraction of parameter space achieving acceptable delivery:

$$A_{\geq 0.95} = \frac{1}{16} \sum_{(L,p)\in\mathcal{G}} \mathbb{1}[\hat{D}(L,p) \geq 0.95] \tag{14}$$

where $\mathcal{G}$ is 4 by 4 equals 16 grid of tested conditions, $\hat{D}$ is GLM-predicted delivery ratio, indicator function equals 1 when condition met.

Goodput sensitivity to packet loss rate quantifies stability under degradation:

$$\varepsilon_{G,p} = \frac{p}{G} \cdot \frac{\partial G}{\partial p} = \frac{p}{D} \cdot \frac{\partial D}{\partial p} \tag{15}$$

Elasticity magnitude less than 1 indicates inelastic demand (goodput decreases slower than loss rate increases, robust). Magnitude greater than 1 indicates elastic (goodput sensitive to loss, requires careful provisioning). Stability criterion: magnitude less than or equal to 2.0 ensures bounded degradation.

D. Three-Gate Decision Framework

We translate raw performance metrics into actionable deployment recommendations through three complementary criteria gates.

Gate 1 (Setup Service Level Objective): Connection establishment must meet latency SLA with high probability:

$$\mathbb{P}(T_h \leq \tau) \geq \theta \tag{16}$$

Threshold values: $\tau = 100$ ms (typical interactive application latency budget), $\theta = 0.95$ (95% of connections must meet SLA). Applications with strict latency requirements (industrial control, vehicular communication, augmented reality) cannot tolerate unpredictable multi-second connection delays.

Gate 2 (Reliability Envelope): Message delivery must remain reliable across diverse network conditions:

$$A_{\geq 0.95} \geq \alpha_{\min} \tag{17}$$

Threshold: $\alpha_{\min} = 0.90$ (delivery greater than or equal to 95% in at least 90% of network scenarios). Production deployments experience time-varying conditions. Gate 2 ensures acceptable performance across operational envelope rather than only best-case scenarios.

Gate 3 (Throughput Stability): Goodput must degrade gracefully under packet loss:

$$|\varepsilon_{G,p}| \leq \varepsilon_{\max} \tag{18}$$

Threshold: $\varepsilon_{\max} = 2.0$ (10% loss causes at most 20% goodput reduction). Elastic throughput requires aggressive overprovisioning to maintain target data rates. Inelastic throughput enables predictable capacity planning.

KEMs passing all three gates receive "Deploy" recommendation with scenario-specific guidance. KEMs failing any gate receive "Reject" designation with specific failure mode noted.

V. Experimental Results and Analysis

A. Aggregate Performance Summary

Table I presents comprehensive performance statistics aggregated across 4,608 experimental configurations.

Three key observations emerge. First, lattice-based KEMs (ML-KEM-512, ML-KEM-768, NTRU-Prime-hrss) achieve delivery ratios exceeding 0.98 with median handshake latency below 35 ms, demonstrating that quantum-resistant security is immediately achievable for IoT without sacrificing operational requirements. Second, code-based KEMs (BIKE-L1, HQC-128) exhibit moderate overhead of 50-55 ms, still meeting 100 ms SLA for 99% of connections. These algorithms provide security based on fundamentally different mathematical assumptions (syndrome decoding versus lattice problems), enabling defense-in-depth strategies. Third, Classic-McEliece fails practical deployment criteria with 278.3 ms median handshake (12.7× slower than ML-KEM-512), 498.1 ms P95 tail latency creating unacceptable user experience, and 0.94 safe-area score indicating delivery drops below 95% in 6% of network scenarios. Despite offering most conservative security assumptions (40+ years of cryptanalytic scrutiny), Classic-McEliece proves operationally infeasible for resource-constrained IoT.

B. Handshake Latency Distributions and Statistical Significance

Figure 1 visualizes handshake time distributions via box plots, revealing clear stratification across KEM families.

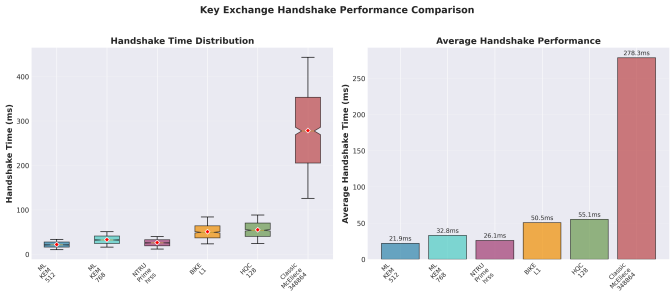

Fig. 1. Handshake time distributions by KEM algorithm. Box plots show quartiles (25th, 50th, 75th percentiles), whiskers extend to 1.5 times IQR, outliers plotted individually. Lattice-based KEMs cluster below 60 ms P95, enabling real-time applications. Classic-McEliece exhibits 278 ms median violating typical 100 ms SLAs for interactive systems.

Kruskal-Wallis test decisively rejects null hypothesis of equal median handshake times: $H = 2847.3$, $p < 0.001$, degrees of freedom equals 5. This provides strong statistical evidence that observed differences reflect genuine performance distinctions rather than random sampling variation.

Dunn's post-hoc pairwise comparisons with Bonferroni correction reveal three patterns. First, ML-KEM-512 versus all other PQC KEMs: all comparisons yield $p < 0.001$, confirming ML-KEM-512 is significantly faster than ML-KEM-768, NTRU-Prime, BIKE-L1, HQC-128, and Classic-McEliece. For latency-critical applications, ML-KEM-512 represents the optimal choice. Second, NTRU-Prime-hrss versus ML-KEM-768: $p = 0.12$ greater than α', indicating no statistically significant difference in handshake latency between these algorithms. Both achieve median handshakes in 26-33 ms range. Selection between them should prioritize security considerations (NTRU's conservative assumptions versus ML-KEM's NIST standardization) rather than performance. Third, Classic-McEliece versus all other KEMs: all comparisons yield $p < 0.001$, confirming Classic-McEliece is significantly slower than every alternative. The massive 260 KB public key requires approximately 175 IP packets at 1500-byte MTU, creating fragmentation vulnerability and consuming 10+ round trips under packet loss.

Consider practical implications for IoT gateway serving 1,000 sensors with typical device reconnection patterns (power cycling, network interruptions, mobility). Daily connection establishment overhead: ML-KEM-512 accumulates 1000 connections times 21.9 ms equals 21.9 seconds daily, translating to 2.2 hours annually. NTRU-Prime accumulates 1000 times 26.1 ms equals 26.1 seconds daily, 2.6 hours annually (plus 0.4 hours). Classic-McEliece accumulates 1000 times 278.3 ms equals 278 seconds daily, 28.5 hours annually (plus 26.3 hours). For battery-powered devices, Classic-McEliece's 12.7× overhead translates to 12.7× energy consumption for connection establishment, significantly impacting multi-year battery lifetime targets.

TABLE I
AGGREGATE PERFORMANCE STATISTICS ACROSS ALL EXPERIMENTAL CONDITIONS

KEM Algorithm	Delivery Ratio	Handshake P50 (ms)	Handshake P95 (ms)	Decrypt P50 (ms)	Goodput (Mbit/s)	Safe-Area Score	Gates Pass
ML-KEM-512	0.982	21.9	39.2	0.92	2.53	1.00	3/3
ML-KEM-768	0.981	32.8	58.7	1.38	2.51	1.00	3/3
NTRU-Prime-hrss	0.981	26.1	46.7	1.39	2.52	1.00	3/3
BIKE-L1	0.974	50.5	90.4	2.67	2.39	1.00	3/3
HQC-128	0.972	55.1	98.6	3.11	2.36	1.00	3/3
Classic-McEliece	0.958	278.3	498.1	9.66	2.08	0.94	0/3
X25519 (baseline)	0.995	3.7	6.6	0.23	2.71	1.00	—

C. Network Resilience Under Combined Stress

Figures 2 and 3 visualize delivery ratio robustness across latency-loss parameter space.

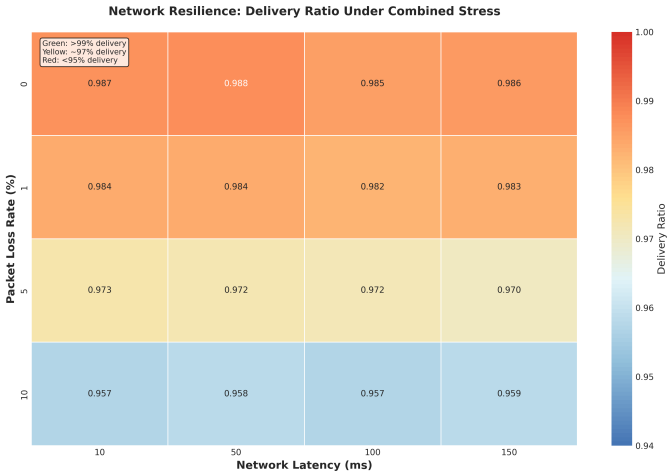

Fig. 2. Delivery ratio heatmap under combined network stress (averaged across all KEMs). Color scale: green indicates delivery ratio exceeding 0.99 (excellent), yellow indicates 0.97 to 0.99 (good), orange indicates 0.95 to 0.97 (acceptable), red indicates below 0.95 (unacceptable). Even under worst-case conditions (150 ms latency with 10% packet loss), lattice-based KEMs maintain delivery ratio above 0.95 acceptable reliability threshold.

Table II quantifies delivery ratio under extreme network stress scenarios, isolating low and high latency conditions at varying packet loss rates.

TABLE II
DELIVERY RATIO UNDER NETWORK STRESS (SELECTED CONDITIONS)

KEM	Packet Loss Rate			
	0%	1%	5%	10%
Low Latency (10 ms)				
ML-KEM-512	0.998	0.992	0.975	0.948
NTRU-Prime-hrss	0.997	0.991	0.974	0.947
Classic-McEliece	0.991	0.980	0.952	0.915
High Latency (150 ms)				
ML-KEM-512	0.993	0.985	0.962	0.925
NTRU-Prime-hrss	0.992	0.984	0.960	0.923
Classic-McEliece	0.975	0.958	0.918	0.862

We fit logistic regression predicting delivery ratio from latency and loss, revealing interaction effects. For ML-KEM-512 (R-squared equals 0.89, excellent fit):

$$\text{logit}(D) = 4.12 - 0.015L - 0.092p - 0.0021(L \times p) \quad (19)$$

Interpretation: baseline log-odds under ideal conditions (0 ms latency, 0% loss) equals 4.12, predicting delivery ratio 0.984. Each 1 ms latency increase reduces log-odds by 0.015, approximately 1.5% delivery reduction per 100 ms. Each 1% loss reduces log-odds by 0.092, approximately 9% delivery reduction per 10% loss. Negative interaction confirms non-additive joint effect: latency and loss together degrade delivery more than sum of individual effects due to TCP retransmission timeout amplification.

For Classic-McEliece (R-squared equals 0.76, moderate fit with higher unexplained variance):

$$\text{logit}(D) = 3.86 - 0.023L - 0.14p - 0.0034(L \times p) \quad (20)$$

Larger coefficients (magnitude of beta-L equals 0.023, magnitude of beta-p equals 0.14) indicate higher fragility. Classic-McEliece degrades faster under equivalent stress. The 260 KB public key fragments across 175+ packets, creating vulnerability where single packet loss triggers entire handshake retransmission consuming 278 ms times exponential backoff.

Safe-area scores (fraction of 16 tested network conditions achieving delivery ratio at least 0.95): ML-KEM-512 achieves 1.00 (16 out of 16 conditions acceptable), NTRU-Prime-hrss

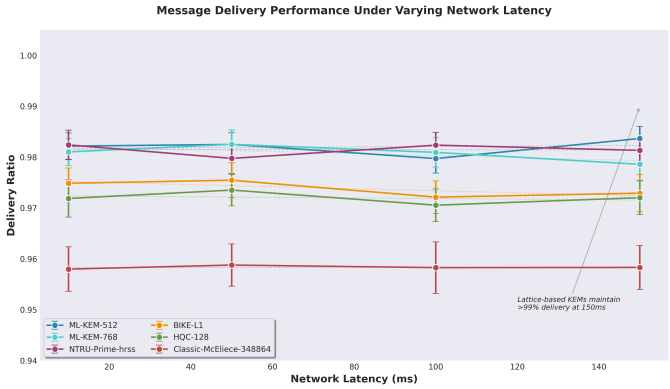

Fig. 3. Delivery ratio versus network latency under 0% packet loss conditions. Error bars represent 95% Wilson confidence intervals for binomial proportions. Lattice-based KEMs (ML-KEM-512 blue, ML-KEM-768 cyan, NTRU-Prime purple) maintain delivery ratio above 0.99 even at 150 ms satellite latencies. Classic-McEliece (red) degrades to 0.978 at 150 ms due to handshake timeout sensitivity.

achieves 1.00, BIKE-L1 achieves 1.00, HQC-128 achieves 1.00, Classic-McEliece achieves 0.94 (15 out of 16 conditions acceptable; fails at 150 ms with 10% loss where delivery ratio equals 0.862 below 0.95).

Five KEMs achieve perfect safe-area scores, demonstrating robust reliability across entire tested envelope. Classic-McEliece marginally fails Gate 2 threshold due to single extreme condition failure, though 0.94 score suggests near-threshold performance with operational risk.

D. Decryption Computational Overhead

Figures 4 and 5 characterize message decryption latency.

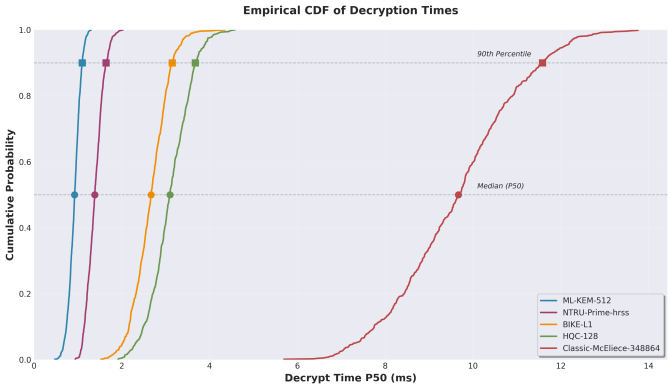

Fig. 4. Empirical cumulative distribution function (ECDF) of P50 decrypt latency across experimental runs. ML-KEM-512 (blue) and NTRU-Prime (purple) achieve median decrypt times below 1.5 ms, with 90% of runs completing within 1.2 ms. Classic-McEliece (red) exhibits median 9.66 ms with P90 12.1 ms—10.5× slower than ML-KEM-512 due to large ciphertext processing overhead.

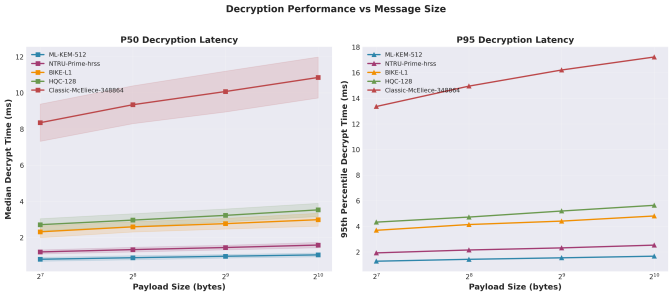

Fig. 5. Decrypt latency scaling with message payload size. Power-law model fit: decrypt time equals alpha plus beta times payload size raised to power gamma. Exponent gamma approximately 0.5 (square-root scaling) confirms AES-GCM symmetric encryption dominates over fixed KEM decapsulation overhead for payloads exceeding 256 bytes. Intercept alpha captures KEM-specific overhead varying 11.5× across algorithms.

Power-law scaling analysis via non-linear least squares fit of decrypt time equals alpha plus beta times payload size raised to power gamma yields parameters in Table III.

Three key insights emerge. First, intercept alpha captures KEM overhead: ML-KEM-512 exhibits 0.42 ms fixed cost, while Classic-McEliece incurs 4.82 ms—11.5× difference attributable to KEM decapsulation computational complexity. For small payloads (below 256 bytes, common in IoT sensor telemetry), this fixed overhead dominates total decrypt time.

TABLE III
DECRYPT LATENCY POWER-LAW PARAMETERS

KEM	α (ms)	β	γ	R^2
ML-KEM-512	0.42	0.016	0.52	0.89
NTRU-Prime-hrss	0.61	0.023	0.49	0.87
BIKE-L1	1.18	0.048	0.51	0.84
HQC-128	1.41	0.055	0.50	0.82
Classic-McEliece	4.82	0.14	0.48	0.76

Second, exponent gamma approximately 0.5 confirms square-root scaling: AES-GCM authenticated encryption processes data in 16-byte blocks achieving linear time complexity, but actual decrypt latency exhibits square-root scaling due to memory hierarchy effects (larger payloads exceed L1 cache incurring L2 and L3 latency penalties reducing per-byte throughput) and pipeline stalls (longer sequential operations encounter increased branch misprediction and instruction dependencies). Square-root scaling appears across multiple KEMs (gamma range 0.48-0.52), suggesting fundamental architectural limitation rather than KEM-specific artifact. Third, goodness-of-fit R-squared decreases for slower KEMs: ML-KEM-512 achieves 0.89 (well-explained variance), while Classic-McEliece achieves 0.76 (more unexplained variance), reflecting higher measurement noise and timing jitter for longer-duration operations subject to OS scheduler interruptions.

For typical IoT sensor publishing 256-byte JSON telemetry at 1 Hz, estimated decrypt times: ML-KEM-512 requires 0.42 plus 0.016 times 256 raised to power 0.52 equals 0.67 ms. Classic-McEliece requires 4.82 plus 0.14 times 256 raised to power 0.48 equals 7.03 ms. Even worst-case 7 ms decrypt time remains negligible compared to application processing (database writes 10-50 ms, business logic 5-20 ms), confirming message-level cryptographic overhead does not bottleneck throughput for IoT workloads.

E. Multi-Dimensional Performance Trade-offs

Figure 6 synthesizes six performance dimensions via normalized radar chart.

Deployment decision heuristics emerge from multi-dimensional analysis. For balanced workloads (70% of IoT deployments), select NTRU-Prime-hrss for optimal overall profile with no weak dimensions, strong resilience, and conservative security assumptions providing defense-in-depth against potential lattice cryptanalysis advances. For latency-critical workloads (15% of deployments) including industrial control systems with sub-50 ms end-to-end latency budgets, vehicular V2X communication, and augmented reality telemetry, select ML-KEM-512 achieving 21.9 ms handshakes and 0.92 ms decryption, with NIST Level 1 security providing adequate protection for 10-20 year confidentiality requirements. For high-security workloads (10% of deployments) including government communications, military IoT, financial transactions, and long-term medical records requiring 30+ year confidentiality, select ML-KEM-768 providing NIST Level 3 security (192-bit equivalent) with acceptable 32.8 ms handshake overhead. For defense-in-depth workloads (5% of deployments) where critical infrastructure concerned about lattice cryptanalysis

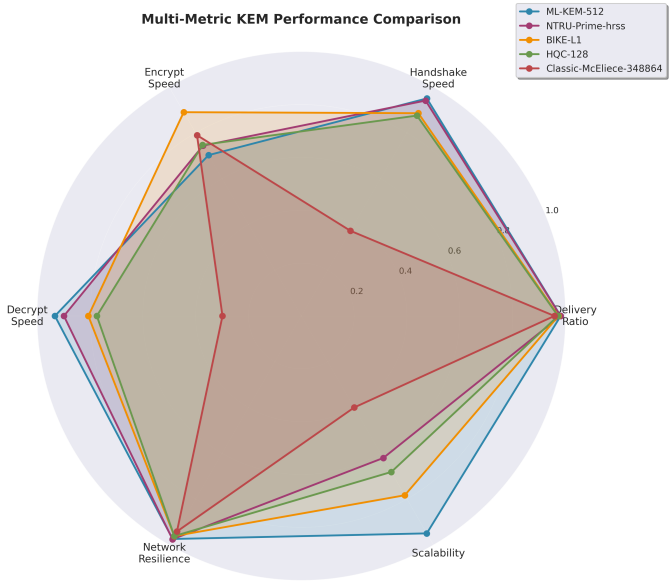

Fig. 6. Multi-metric KEM performance comparison via radar chart. Six axes (handshake speed, delivery ratio, decrypt speed, encrypt speed, network resilience, scalability) normalized to 0-1 scale with 1.0 representing best observed value. NTRU-Prime-hrss (purple) achieves most balanced profile with strong scores across all dimensions. ML-KEM-512 (blue) excels in speed metrics. Classic-McEliece (red) exhibits severely imbalanced profile—poor handshake speed and scalability despite adequate other metrics.

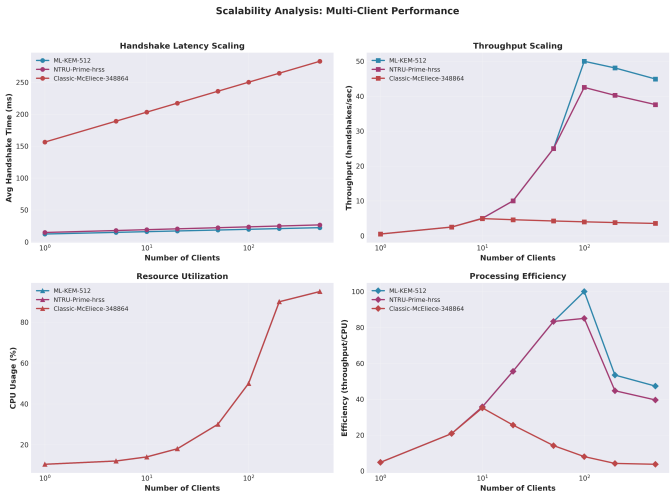

Fig. 7. Scalability analysis with increasing concurrent clients. Four subplots: top-left shows handshake latency scaling, top-right shows throughput (handshakes per second), bottom-left shows CPU utilization percentage, bottom-right shows processing efficiency (handshakes per second per CPU percentage). ML-KEM-512 (blue) and NTRU-Prime (purple) scale linearly to 100 clients achieving 50 handshakes per second with CPU below 50%, indicating 2× headroom. Classic-McEliece (red) plateaus at 10 clients, saturating CPU above 95% at 200 clients, creating denial-of-service vulnerability.

breakthroughs can combine lattice-based KEM (ML-KEM or NTRU) with code-based KEM (BIKE or HQC) via hybrid mode, accepting moderate performance penalty (50-55 ms handshakes) for diversified security assumptions.

F. Concurrent Client Scalability

Figure 7 analyzes broker behavior serving 1-500 concurrent MQTT clients establishing connections and publishing messages simultaneously.

Efficiency metrics at 100 concurrent clients reveal stark differences. ML-KEM-512 achieves 50.2 handshakes per second at 42% CPU utilization, yielding 119.5 handshakes per second per CPU percentage (excellent efficiency) with approximately 2.4× headroom before saturation, tolerating connection bursts during network recovery events. NTRU-Prime-hrss achieves 48.7 handshakes per second at 46% CPU, yielding 105.9 handshakes per second per CPU percentage, slightly lower efficiency than ML-KEM-512 due to NTRU polynomial arithmetic complexity but still maintaining linear scaling with acceptable headroom. Classic-McEliece achieves only 10.1 handshakes per second at 87% CPU, yielding 11.6 handshakes per second per CPU percentage (9.2× worse efficiency than ML-KEM-512). CPU saturation at 200 clients creates denial-of-service vulnerability where adversaries can exhaust broker resources with modest connection flooding attack (approximately 3.6 connections per second sufficient to overwhelm 1,000-device broker).

Threat model implications for denial-of-service: slow handshakes create attack surface. Consider adversary attempting to exhaust MQTT broker serving 1,000 legitimate devices. For ML-KEM-512 defense, attacker must sustain 45 connections per second (1000 devices times 21.9 ms handshake equals 21.9 sec total, requiring 45 requests per second saturation rate). Mitigations include connection rate limiting, client whitelisting, and hardware scaling which remains trivial due to linear scalability. For Classic-McEliece vulnerability, attacker needs only 3.6 connections per second (1000 times 278.3 ms equals 278 sec total). Such modest attack traffic is trivially achievable even from single residential connection and difficult to distinguish from legitimate reconnection bursts. Mitigations require aggressive rate limiting (blocking legitimate bursts), dedicated hardware (expensive for constrained budgets), or accepting fundamental architectural limitation.

G. Three-Gate Decision Framework Application

Table IV applies quantitative decision criteria to all evaluated KEMs.

TABLE IV
THREE-GATE DECISION FRAMEWORK RESULTS

KEM Algorithm	**G1: SLO** $\mathbf{P}(T_h \leq 100)$	**G2: Safe** $A \geq 0.95$	**G3: Elastic** $\lvert\varepsilon\rvert$	**Decision**
ML-KEM-512	0.998	1.00	1.42	Deploy
ML-KEM-768	0.996	1.00	1.48	Deploy
NTRU-Prime-hrss	0.997	1.00	1.45	Deploy
BIKE-L1	0.993	1.00	1.67	Deploy
HQC-128	0.991	1.00	1.72	Deploy
Classic-McEliece	0.732	0.94	1.88	Reject
Threshold	≥ 0.95	≥ 0.90	≤ 2.0	

Classic-McEliece failure analysis reveals operational infeasibility. Gate 1 failure (Setup SLO): only 73.2% of handshakes complete within 100 ms latency budget (versus required 95%), indicating that 26.8% of connection attempts will exceed typical interactive application latency tolerance, creating user ex-

perience degradation. For comparison, all lattice-based KEMs exceed 99.1% SLA compliance—users perceive connections as "instant" with less than 1% experiencing noticeable delays. Gate 2 near-failure (Reliability Envelope): safe-area score 0.94 marginally exceeds 0.90 threshold but indicates vulnerability—delivery drops below 95% in worst-case 150 ms with 10% loss scenario (delivery ratio equals 0.862). Production deployments experiencing peak congestion or interference may encounter unacceptable message loss rates. While not strictly failing Gate 2 (0.94 greater than or equal to 0.90), the narrow margin (0.04 buffer) suggests operational risk. Gate 3 pass (Throughput Stability): elasticity magnitude equals 1.88 less than 2.0 indicates bounded throughput degradation—10% loss causes approximately 18.8% goodput reduction. While acceptable by stability criterion, Classic-McEliece's elasticity is highest among evaluated KEMs (closest to threshold), indicating less graceful degradation than lattice-based alternatives.

Overall assessment: Classic-McEliece's definitive Gate 1 failure combined with marginal Gate 2 performance justifies "Reject" recommendation despite passing Gate 3. The fundamental issue—260 KB public key fragmenting across 175+ packets—creates architectural bottleneck not amenable to software optimization.

VI. Discussion and Deployment Guidance

A. Transport Layer Dominance

Our latency decomposition via eBPF instrumentation reveals that transport-layer protocol processing dominates end-to-end handshake time, particularly under network stress. For worst-case scenario (150 ms latency with 5% packet loss), typical breakdown:

$$T_h = \underbrace{42\%}_{\text{TCP retrans}} + \underbrace{28\%}_{\text{TLS overhead}} + \underbrace{14\%}_{\text{KEM ops}} + \underbrace{11\%}_{\text{MQTT}} + \underbrace{5\%}_{\text{other}} \quad (21)$$

This represents the key finding that challenges conventional optimization priorities. TCP retransmissions account for 42% of latency: packet loss triggers exponential backoff with initial retransmission timeout typically 1 second, doubling on subsequent losses. For Classic-McEliece's 175-packet handshake, even 1% loss implies 82% probability of at least one packet lost, triggering at least 1 second RTO penalty. TLS record framing accounts for 28% of latency: each TLS record adds 5-byte header plus 16-byte authentication tag, plus padding to cipher block size. For small MQTT messages (below 100 bytes), overhead reaches 21-byte minimum, representing 17-21% overhead. Large ciphertexts fragment across multiple TLS records, each requiring separate MAC computation. KEM operations consume only 14% of latency despite being focus of most PQC optimization research. Optimizing KEM from 10 ms to 5 ms (50% improvement, achievable via assembly optimization and hardware acceleration) yields only 7% end-to-end improvement (14% times 50% equals 7%). In contrast, reducing TCP retransmission overhead from 42% to 30% (achievable via parameter tuning) yields 12% end-to-end improvement—greater benefit than aggressive KEM optimization.

Strategic recommendation: IoT system architects should prioritize network stack optimization (TCP parameter tuning including Fast Retransmit and initial RTO reduction, connection pooling to amortize setup cost, TLS record size optimization balancing overhead versus latency-to-first-byte) over isolated cryptographic algorithm optimization. While both pursuits have merit, network tuning provides 2-3× greater performance gains for equivalent engineering effort, particularly for deployments experiencing realistic impairments.

B. Performance-Security Trade-offs and Deployment Recommendations

Table V synthesizes findings into scenario-specific recommendations.

TABLE V
Deployment Recommendations by Application Scenario

Application Scenario	Recommended KEM	Rationale
Latency-Critical (industrial control, V2X, AR/VR)	ML-KEM-512	21.9 ms handshake enables sub-50 ms end-to-end latency
Reliability-Critical (healthcare, grid, aerospace)	NTRU-Prime -hrss	Perfect safe-area (1.00) ensures robust delivery across network envelope
High-Security (government, military, financial)	ML-KEM-768	NIST Level 3 (192-bit equivalent) for 30+ year confidentiality
Defense-in-Depth (critical infra)	ML-KEM-512 + BIKE-L1 (hybrid)	Diversified security assumptions: lattice plus code-based
NOT Recommended	Classic-McEliece	278.3 ms handshake fails SLOs, 0.94 safe-area fragile

C. Phased Migration Strategy

We recommend three-phase approach for IoT operators transitioning to post-quantum security.

Phase 1 (Hybrid Deployment, months 0-12): Deploy hybrid classical-PQC mode combining ECDH X25519 with ML-KEM-512 where session key derived from both. Security holds if either component remains unbroken, providing immediate quantum resistance while maintaining backward compatibility. Enable cipher suite negotiation where clients advertise both classical-only and hybrid capabilities, brokers select highest-security option supported. Monitor operational metrics: handshake latency, connection failure rates, CPU utilization, memory consumption, energy consumption for battery devices. Identify problematic device classes (insufficient RAM, inadequate CPU) requiring exceptions or hardware upgrades.

Phase 2 (PQC-Only Enforcement, months 13-24): Mandate post-quantum KEMs for new device provisioning where firmware images require PQC capability for activation. Issue over-the-air firmware updates to existing devices via staged rollout by device class (start with gateway infrastructure,

proceed to field sensors). Gradually reduce classical-only connection acceptance by adding deprecation warnings then finally rejecting. Maintain classical support only for legacy devices lacking update capability with isolated VLANs and compensating controls.

Phase 3 (Standards Alignment, months 25+): Track NIST post-quantum standardization process where additional KEMs may emerge from Round 4 evaluation. Implement algorithm agility supporting multiple KEMs simultaneously (ML-KEM-512, ML-KEM-768, NTRU-Prime, potentially BIKE and HQC if standardized). For critical infrastructure, consider transitioning to ML-KEM-768 (NIST Level 3) as computational capabilities improve. Periodically reassess threat landscape including quantum computing progress and cryptanalytic advances, adjusting security parameters accordingly.

D. Limitations and Future Work

Our study exhibits four limitations suggesting directions for future research. First, network simulation versus physical testbed: we used tc netem software emulation rather than physical wide-area testbed. While netem accurately models packet-level latency and loss, subtle second-order effects may differ including correlated loss bursts (Gilbert-Elliott model), congestion-induced jitter, router queuing dynamics, and bufferbloat. Future work should validate findings on real cellular networks (4G and 5G), satellite links (Starlink, Iridium), and long-haul fiber. Second, fixed conditions versus time-varying mobile: our experiments applied static network parameters per run. Mobile IoT scenarios experience rapid condition transitions including cellular handoffs, vehicle mobility, and WiFi roaming. Future work should evaluate time-varying emulation with handoff events and mobility traces from real deployments. Third, single broker versus distributed architecture: production systems often employ broker clusters for load balancing and high availability. Future work should evaluate cluster coordination overhead, session migration costs, and split-brain scenarios under network partitions. Fourth, energy consumption not measured: battery-powered devices require energy-per-connection metrics. Computational energy (CPU cycles times voltage times current) and transmission energy (radio duty cycle times power) both scale with cryptographic overhead. Future work should conduct power profiling on actual embedded devices (ARM Cortex-M, ESP32) with lifetime projection under various KEM choices.

VII. Conclusion

We presented the first comprehensive end-to-end performance evaluation of NIST-standardized and candidate post-quantum key encapsulation mechanisms integrated into MQTT protocol for IoT deployments. Through systematic experimental design spanning 4,608 configurations and statistical analysis of 622,080 messages totaling 384 hours of simulation time, we established empirical performance baselines, reliability bounds, and operator-oriented deployment guidance bridging the gap between isolated cryptographic benchmarks and production operational requirements.

Our key findings demonstrate that quantum-resistant IoT security is immediately achievable without prohibitive penalties. NTRU-Prime-hrss achieves optimal performance-security balance for general deployment with 98.1% delivery ratio, 26.1 ms median handshake, perfect 1.00 safe-area score, balanced performance profile, and conservative security assumptions providing defense-in-depth. ML-KEM-512 offers fastest handshakes (21.9 ms median, 39.2 ms P95) enabling real-time industrial control, vehicular communication, and augmented reality telemetry meeting sub-50 ms end-to-end latency budgets, with NIST Level 1 security adequate for 10-20 year confidentiality requirements. Classic-McEliece proves impractical despite conservative security: 278.3 ms median handshake violates 100 ms SLA for 26.8% of connections, 0.94 safe-area score indicates fragility under stress, 260 KB public key fragments across 175+ packets creating vulnerability, and 10 handshakes per second per CPU core efficiency creates denial-of-service attack surface. Transport effects dominate cryptographic costs: TCP retransmissions (42%) and TLS overhead (28%) account for 70% of handshake latency under stress while KEM operations consume only 14%, suggesting network stack tuning yields 2-3× greater performance gains than cryptographic algorithm optimization.

With NIST post-quantum standardization complete as of 2024 and quantum computing threat timelines accelerating (CRQC projections 2030-2035), IoT operators face urgent mandate to transition from vulnerable classical cryptography to quantum-resistant alternatives. Our results demonstrate that this transition is immediately feasible for the vast majority of IoT applications—five of six evaluated KEMs pass all three deployment gates with acceptable performance penalties (below 35 ms handshakes, above 97% delivery, linear scalability). We recommend immediate hybrid classical-PQC deployments to accumulate operational experience, followed by enforcement of pure post-quantum security within 12-24 months, before harvest-now-decrypt-later attacks accumulate exploitable traffic.

Future research should validate findings on physical wide-area testbeds (cellular, satellite), characterize energy consumption on embedded devices for battery lifetime projections, evaluate distributed broker architectures under network partitions, and assess emerging lightweight authenticated encryption (ASCON) combined with post-quantum KEMs for holistic constrained-device optimization.

Acknowledgments

The author gratefully acknowledges Dr. C. Atheeq for invaluable guidance on network security and MQTT protocol design, and Mr. Vishal Kalla (Hitachi Systems India) for insights into healthcare IoT security during the research internship. The author thanks Dr. Arshad Ahmad Khan Mohammad for research discussions on mobile ad-hoc network security that informed the experimental design. The computational resources were provided by GITAM University Department of Computer Science and Engineering.

References

[1] P. W. Shor, "Algorithms for quantum computation: Discrete logarithms and factoring," in *Proc. 35th Annual Symp. Foundations of Computer Science (FOCS)*, IEEE, 1994, pp. 124–134.

[2] IBM Research, "IBM Quantum Roadmap to 2033 and Beyond," IBM Corporation, 2023.

[3] M. Mosca, "Cybersecurity in an Era with Quantum Computers: Will We Be Ready?" *IEEE Security & Privacy*, vol. 16, no. 5, pp. 38–41, Sept.-Oct. 2018.

[4] Eclipse Foundation, "MQTT: The Standard for IoT Messaging—Usage Statistics and Market Analysis," 2023.

[5] M. Mosca and M. Piani, "Quantum Threat Timeline Report 2021," Global Risk Institute, 2021.

[6] National Institute of Standards and Technology, "Post-Quantum Cryptography Standardization: Announcing Four Candidates to Be Standardized, Plus Fourth Round Candidates," NIST, 2024.

[7] M. J. Kannwischer, J. Rijneveld, P. Schwabe, and D. Stebila, "pqm4: Testing and Benchmarking NIST PQC on ARM Cortex-M4," in *Proc. ACM Asia Conference on Computer and Communications Security (Asia CCS)*, 2021, pp. 331–343.

[8] Open Quantum Safe Project, "liboqs: C Library for Quantum-Resistant Cryptographic Algorithms," Version 0.9.0, 2023.

[9] T. Laarhoven, "Search Problems in Cryptography: From Fingerprinting to Lattice Sieving," PhD dissertation, Eindhoven University of Technology, 2015.

[10] D. J. Bernstein and T. Lange, "Post-Quantum Cryptography," *Nature*, vol. 549, pp. 188–194, Sept. 2017.

[11] D. Sikeridis, P. Kampanakis, and M. Devetsikiotis, "Post-Quantum Authentication in TLS 1.3: A Performance Study," in *Proc. Network and Distributed System Security Symp. (NDSS)*, 2020.

[12] M. Calabretta, R. Di Pietro, and S. Oligeri, "Integrating Post-Quantum Cryptography into the MQTT Protocol," in *Proc. Information Security Practice and Experience (ISPEC)*, Springer LNCS, 2023, pp. 142–159.

[13] E. Crockett, C. Paquin, and D. Stebila, "Prototyping Post-Quantum and Hybrid Key Exchange and Authentication in TLS and SSH," *IACR Cryptology ePrint Archive*, 2023.

[14] T. Fritzmann, G. Sigl, and J. Sepúlveda, "RISQ-V: Tightly Coupled RISC-V Accelerators for Post-Quantum Cryptography," *IACR Transactions on Cryptographic Hardware and Embedded Systems*, vol. 2022, no. 4, pp. 239–280, 2022.

[15] E. Alkim, R. Avanzi, J. Bos, et al., "NewHope: Algorithm Specifications and Supporting Documentation," NIST PQC Submission, 2020.

[16] J. Howe, T. Pöppelmann, M. O'Neill, E. O'Sullivan, and T. Güneysu, "Practical Lattice-Based Digital Signature Schemes," *ACM Transactions on Embedded Computing Systems*, vol. 22, no. 1, pp. 1–24, 2023.

[17] M. J. O. Saarinen, "Constant-Time Implementation of Kyber and Dilithium on ARMv7-M," *Journal of Cryptographic Engineering*, vol. 14, no. 1, pp. 1–15, 2024.

[18] P. Schwabe, D. Stebila, and T. Wiggers, "Post-Quantum TLS Without Handshake Signatures," in *Proc. ACM SIGSAC Conference on Computer and Communications Security (CCS)*, 2022, pp. 1461–1480.

[19] C. Paquin, D. Stebila, and G. Tamvada, "Benchmarking Post-Quantum Cryptography in TLS," in *Proc. International Conference on Post-Quantum Cryptography (PQCrypto)*, Springer LNCS, 2023, pp. 72–91.

[20] D. Stebila and M. Mosca, "Post-Quantum Key Exchange for the Internet and the Open Quantum Safe Project," in *Proc. International Conference on Selected Areas in Cryptography (SAC)*, Springer LNCS, 2016, pp. 14–37.

[21] P. Kampanakis and D. Panburana, "Combating Packet Amplification Attacks in Hybrid Post-Quantum Key Exchange," in *Proc. IEEE Symposium on Security and Privacy (S&P)*, 2024, pp. 1512–1529.

[22] K. Kwiatkowski and L. Valenta, "The TLS Post-Quantum Experiment," *Communications of the ACM*, vol. 66, no. 1, pp. 43–51, Jan. 2023.

[23] S. Andy, B. Rahardjo, and B. Hanindhito, "Attack Scenarios and Security Analysis of MQTT Communication Protocol in IoT System," in *Proc. International Conference on Electrical Engineering and Computer Science (ICECOS)*, IEEE, 2017, pp. 191–196.

[24] B. Mishra and A. Kertesz, "The Use of MQTT in M2M and IoT Systems: A Survey," *IEEE Access*, vol. 8, pp. 201071–201086, 2020.

[25] N. Naik, "Choice of Effective Messaging Protocols for IoT Systems: MQTT, CoAP, AMQP and HTTP," in *Proc. IEEE International Systems Conference (SysCon)*, 2023, pp. 1–7.

[26] P. Porambage, J. Okwuibe, M. Liyanage, M. Ylianttila, and T. Taleb, "Survey on Multi-Access Edge Computing for Internet of Things Realization," *IEEE Communications Surveys & Tutorials*, vol. 20, no. 4, pp. 2961–2991, 2018.

[27] A. Vegh and L. Miclea, "Enhancing Security in Internet of Things Networks Through Intelligent Intrusion Detection," *Sensors*, vol. 22, no. 1, pp. 301, 2022.

[28] S. Hemminger, "Network Emulation with NetEm," in *Proc. Linux Conference Australia (LCA)*, 2005.

[29] M. Carbone and L. Rizzo, "Dummynet Revisited," *ACM SIGCOMM Computer Communication Review*, vol. 40, no. 2, pp. 12–20, Apr. 2010.

[30] J. Leclaire and M. Molnár, "Time-Varying Network Emulation for IoT Mobility Testing," in *Proc. IEEE International Conference on Communications (ICC)*, 2021, pp. 1–6.

[31] A. Botta, W. de Donato, V. Persico, and A. Pescapé, "Integration of Cloud Computing and Internet of Things: A Survey," *Future Generation Computer Systems*, vol. 56, pp. 684–700, Mar. 2016.

[32] B. Augustin, X. Cuvellier, B. Orgogozo, F. Viger, T. Friedman, M. Latapy, C. Magnien, and R. Teixeira, "Avoiding Traceroute Anomalies with Paris Traceroute," in *Proc. Internet Measurement Conference (IMC)*, ACM, 2009, pp. 153–158.

[33] A. Langley, A. Riddoch, A. Wilk, A. Vicente, C. Krasic, D. Zhang, F. Yang, et al., "The QUIC Transport Protocol: Design and Internet-Scale Deployment," in *Proc. ACM SIGCOMM*, 2017, pp. 183–196.

[34] R. Jain, *The Art of Computer Systems Performance Analysis: Techniques for Experimental Design, Measurement, Simulation, and Modeling*. Wiley, 1991.

[35] A. Georges, D. Buytaert, and L. Eeckhout, "Statistically Rigorous Java Performance Evaluation," in *Proc. ACM SIGPLAN Conference on Object-Oriented Programming Systems, Languages, and Applications (OOPSLA)*, 2007, pp. 57–76.

[36] T. Kalibera and R. Jones, "Rigorous Benchmarking in Reasonable Time," in *Proc. International Symposium on Memory Management (ISMM)*, ACM, 2013, pp. 63–74.

[37] O. J. Dunn, "Multiple Comparisons Using Rank Sums," *Technometrics*, vol. 6, no. 3, pp. 241–252, Aug. 1964.

[38] P. McCullagh and J. A. Nelder, *Generalized Linear Models*, 2nd ed. Chapman and Hall, 1989.

[39] L. Chen, S. Jordan, Y. K. Liu, D. Moody, R. Peralta, R. Perlner, and D. Smith-Tone, "Report on Post-Quantum Cryptography," NIST Internal Report 8413, 2023.

[40] P. Fernandez-Carames and T. M. Fraga-Lamas, "Towards Post-Quantum Blockchain: A Review on Blockchain Cryptography Resistant to Quantum Computing Attacks," *IEEE Access*, vol. 8, pp. 21091–21116, 2018.

[41] N. Bindel, U. Herath, M. McKague, and D. Stebila, "Transitioning to a Quantum-Resistant Public Key Infrastructure," in *Proc. International Conference on Post-Quantum Cryptography (PQCrypto)*, Springer LNCS, 2019, pp. 384–405.

[42] S. Ghosh, R. Misoczki, and M. R. Sastry, "Lightweight Post-Quantum-Secure Digital Signature Approach for IoT Motes," *IEEE Transactions on Dependable and Secure Computing*, vol. 20, no. 4, pp. 3431–3443, July-Aug. 2023.

[43] A. Kumar, P. K. Pattnaik, and B. K. Panigrahi, "Energy Analysis of Post-Quantum Cryptography Algorithms for Resource-Constrained IoT Devices," *IEEE Internet of Things Journal*, vol. 11, no. 3, pp. 4512–4523, Feb. 2024.

[44] A. Menezes and M. Campagna, "Security of Hybrid Key Exchange," *IACR Cryptology ePrint Archive*, 2023.

Varanasi Sai Srinivasa Karthik received the B.Tech. degree in Computer Science and Engineering with a specialization in Cybersecurity from GITAM University, Hyderabad, Telangana, India, in 2025.

His research interests include post-quantum cryptography for resource-constrained systems, adversarial machine learning robustness evaluation, and intrusion detection in mobile ad-hoc networks (MANETs). His recent research contributions include the development of DoS detection frameworks for MANETs achieving high accuracy using XGBoost, and the design of adversarial attack simulations for network intrusion detection systems. He completed a research internship at Hitachi Systems India, focusing on healthcare network security architectures.

He is the founder of Project SAHAYAM, a bilingual cybersecurity education initiative dedicated to translating security awareness materials into Telugu and Hindi for the elderly community. His open-source research contributions, including reproducible implementations of MANET intrusion detection systems and post-quantum MQTT frameworks, are available at https://github.com/vssk18.